Fuzzy Tolerance Clustering Workbench

User Guide

Aliaksandr Damaratski

2026-11-01

Table of contents

1 About this guide **1**

2 Installation and first start **2**

3 The screen **2**

4 Loading data **3**
4.1 CSV files . . . 3
4.2 Tolerance matrices . . . 4
4.3 The data page . . . 4

5 Choosing an algorithm **4**
5.1 Parameters . . . 7

6 Running an algorithm **9**

7 Reading the results **9**
7.1 Clusters . . . 9
7.2 Membership U . . . 10
7.3 Tolerance . . . 11
7.4 Dendrogram . . . 12
7.5 Validity . . . 12

8 Visualize **13**

9 Export **14**

10 Theory pages **14**
10.1 Fuzzy sets . . . 14
10.2 Inference (FIS) . . . 15

11 Settings **16**

12 Troubleshooting **17**

13 The command-line tool **17**

1 About this guide

The workbench is the desktop application of the *fuzzy-tolerance-clustering* project. It runs the possibilistic clustering algorithms of D. A. Viattchenin's D-AFC family — the algorithms described in the book *Fuzzy Tolerance Clustering: Theory, Algorithms, and Implementations* — on your own data, shows the tolerance relation, the allotments and the validity indices, and turns a clustering solution into a fuzzy inference system.

This guide explains the application screen by screen. The mathematics behind the options is only summarised here; the book is the reference for it. The same algorithms are available without a graphical interface as the

command-line tool `mfuzzy` and as a .NET library (see Section 13).

Source code, data sets and releases: https://github.com/adamaratski/fuzzy-tolerance-clustering.

2 Installation and first start

Requirements. Windows 10 or 11, 64-bit. Nothing else needs to be installed: the archive contains a single executable with the .NET runtime inside.

Installation.

1. Download `FuzzyToleranceClustering-Workbench-<version>-win-x64.zip` from the *Releases* page of the repository. The file `*.zip.sha256` next to it holds the SHA-256 checksum of the archive.
2. Unpack the archive into any folder. It contains `Math.Fuzzy.Desktop.exe`, the folder `datasets` with the benchmark data sets of the book, `LICENSE.txt` and a short `README.txt`.
3. Start `Math.Fuzzy.Desktop.exe`.

The executable is not code-signed. On the first start Windows SmartScreen may report an unknown publisher; choose *More info* and then *Run anyway*.

First start. The application opens the Dashboard with two sample data sets from the `datasets` folder next to the executable: Anderson's Iris data (150 objects, 4 attributes, 3 classes) and the Yang–Wu data (16 objects, 2 attributes). If the folder is missing, the application starts empty and data can be opened with *Open data*.

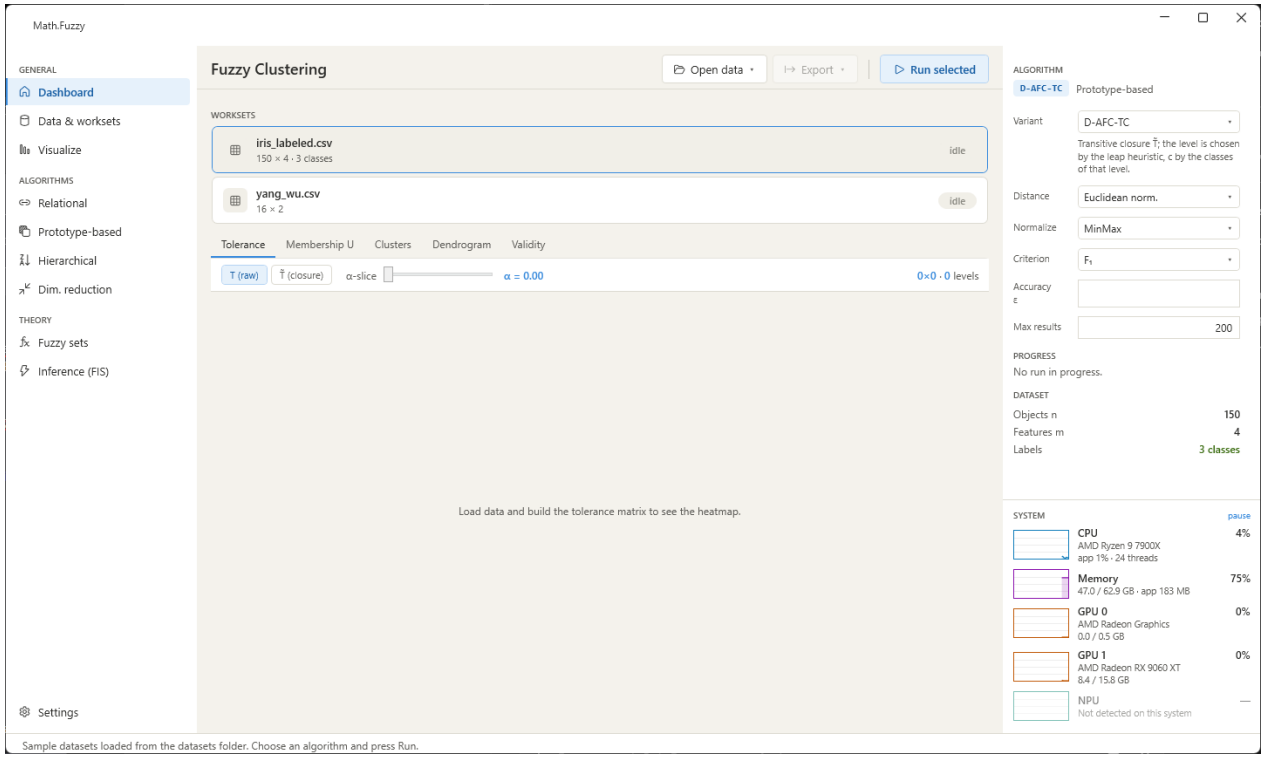


Figure 1: The Dashboard at start-up with the sample data sets.

3 The screen

The window has four areas (see the figure above).

- **Navigation** (left). *General* holds the Dashboard, the data page and the scatter plot. *Algorithms* lists the four categories of algorithms; choosing a category returns to the Dashboard and selects the category's first algorithm. *Theory* holds two pages for exploring fuzzy sets and fuzzy inference.
- **Header** of the Dashboard. *Open data* loads a CSV file or a tolerance matrix, *Export* saves the results of the last run, *Run selected* starts the algorithm configured in the right panel. While a run is in progress

the button turns into *Stop*, next to it the algorithm name and the elapsed time are shown, and a progress bar appears under the header.

- **Workspace** (centre). The list of *worksets* — the data sets opened in this session, each with its status (*idle*, *running…*, *done*, *error*) — and the result tabs *Tolerance*, *Membership U*, *Clusters*, *Dendrogram* and *Validity*. Below the tabs a green bar summarises the best allotment of the last run and three tiles show its number of clusters, purity and adjusted Rand index (the last two only when the data set has class labels).
- **Right panel.** The parameters of the selected algorithm (Section 5), the progress of the current run, a summary of the active data set and the *System* monitor with the load of the processor, the memory and the graphics adapters. The monitor can be paused with the small link in its corner.

The status bar at the bottom reports what happened last: a loaded file, the result of a run with its validity indices, or the reason why a run did not start.

4 Loading data

4.1 CSV files

Open data → *CSV data…* (Ctrl+O) opens a comma-separated file. The rules are:

- A header row is optional. Column names from the header are used as attribute names on the scatter plot and in the inference page.
- Every numeric column is an attribute. Values use a decimal point.
- One column may hold class labels. A text column is treated as the label column automatically; a numeric column is treated as labels when its header is `class`, `label`, `labels`, `type`, `species`, `variety`, `target` or `category`. Labels are never used by the algorithms — they serve only to compute purity, the Rand indices and the mutual information after a run.

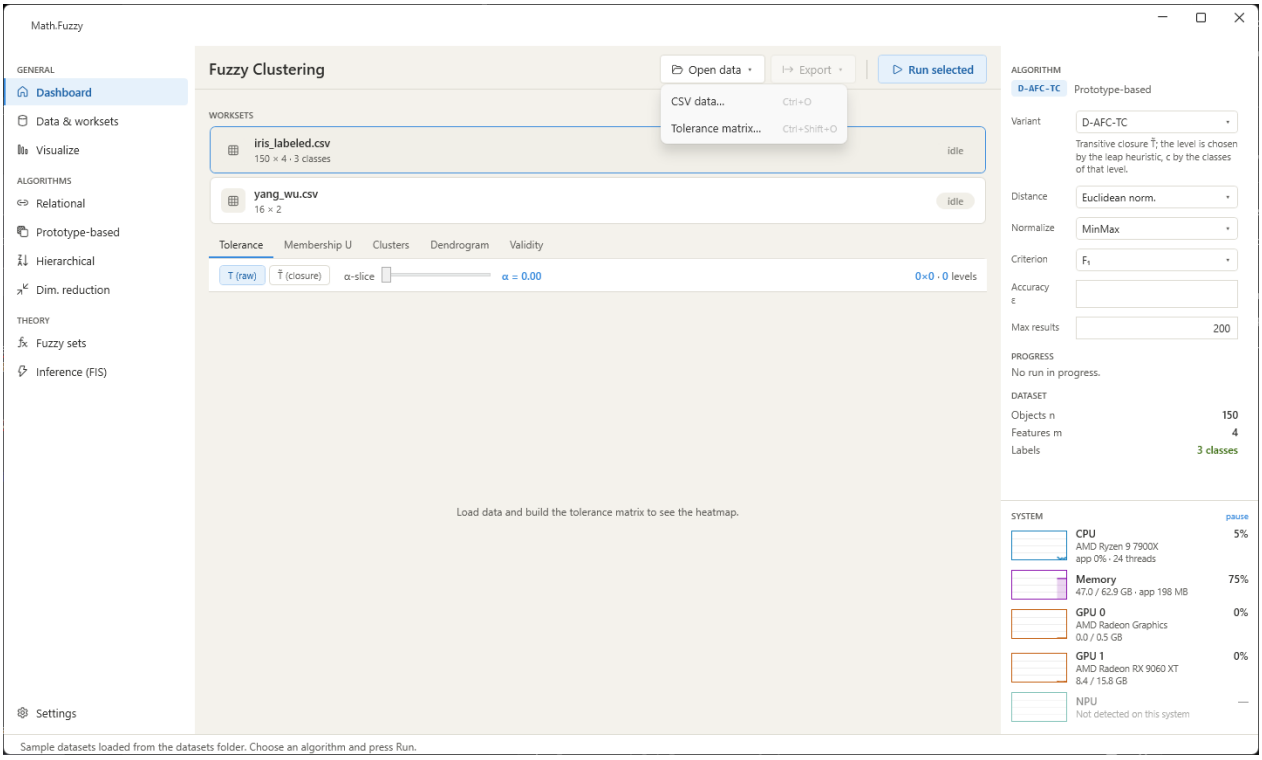


Figure 2: The *Open data* menu.

Every opened file becomes a workset. Click a workset to make it active: the right panel, the result tabs and the scatter plot switch to it, and a run always applies to the active workset. The data sets shipped with the application have the format described above; the file `datasets/README.md` lists their sources.

4.2 Tolerance matrices

Open data → Tolerance matrix… (Ctrl+Shift+O) loads an $n \times n$ matrix T with entries in $[0, 1]$. The diagonal is set to 1 and the matrix is used as the fuzzy tolerance relation directly, without attributes. Such a workset shows the matrix on the *Tolerance* tab and can be examined there, but the algorithms of this version run on attribute data only (Section 12).

4.3 The data page

Data & worksets shows the same list of worksets, the size of the active data set and a preview of its first rows. The buttons *Build T* and *Closure T* compute the tolerance matrix and its max-min transitive closure for the active data set with the distance, normalisation and accuracy of the right panel, and display them on the *Tolerance* tab without running a clustering.

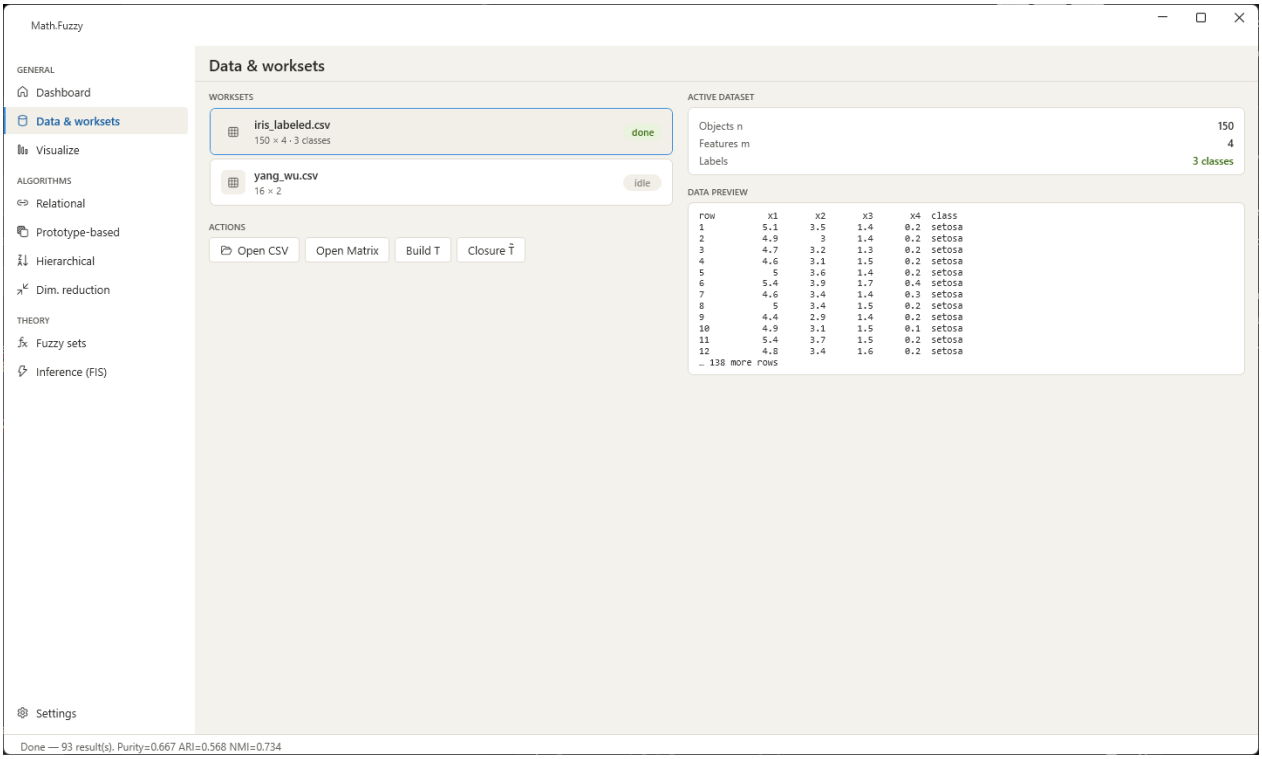


Figure 3: The data page with the preview of the active data set.

5 Choosing an algorithm

The navigation picks a category, the *Variant* box in the right panel picks the algorithm within it, and a short line under the box says what the algorithm does. The blue chip at the top of the panel always shows the selected algorithm.

Table 1: The algorithms of the workbench.

Category	Algorithm	Needs c?	What it does
Relational	D-AFC	yes	The basic algorithm: the best allotment among c clusters over all α-levels of T.
Relational	D-PAFC	no	The principal allotment: the least number of clusters at the first admissible level.

Category	Algorithm	Needs c?	What it does
Prototype-based	D-AFC-TC	no	Replaces T by its transitive closure $\widetilde{T}$; the level is chosen by the leap heuristic, the number of clusters follows from the level.
Prototype-based	D-AFC-TC(α)	no	As D-AFC-TC, but the leap is searched only at levels not below a threshold α^* given by the user.
Prototype-based	D-PAFC-TC	no	The principal allotment on $\widetilde{T}$: its first splitting level.
Prototype-based	D-AFC-TAGA, D-PAFC-TAGA	no	The same with a transitive approximation built by an aggregation operator instead of the closure. Experimental in this version.
Hierarchical	H-AFC-TC	no	The allotments of every level of $\widetilde{T}$, shown as a dendrogram.
Dim. reduction	D-PAFC-DR, D-PAFC-TC-DR, D-PAFC-TAGA-DR	no	Attribute selection: the same algorithms applied to a tolerance relation between the attributes.

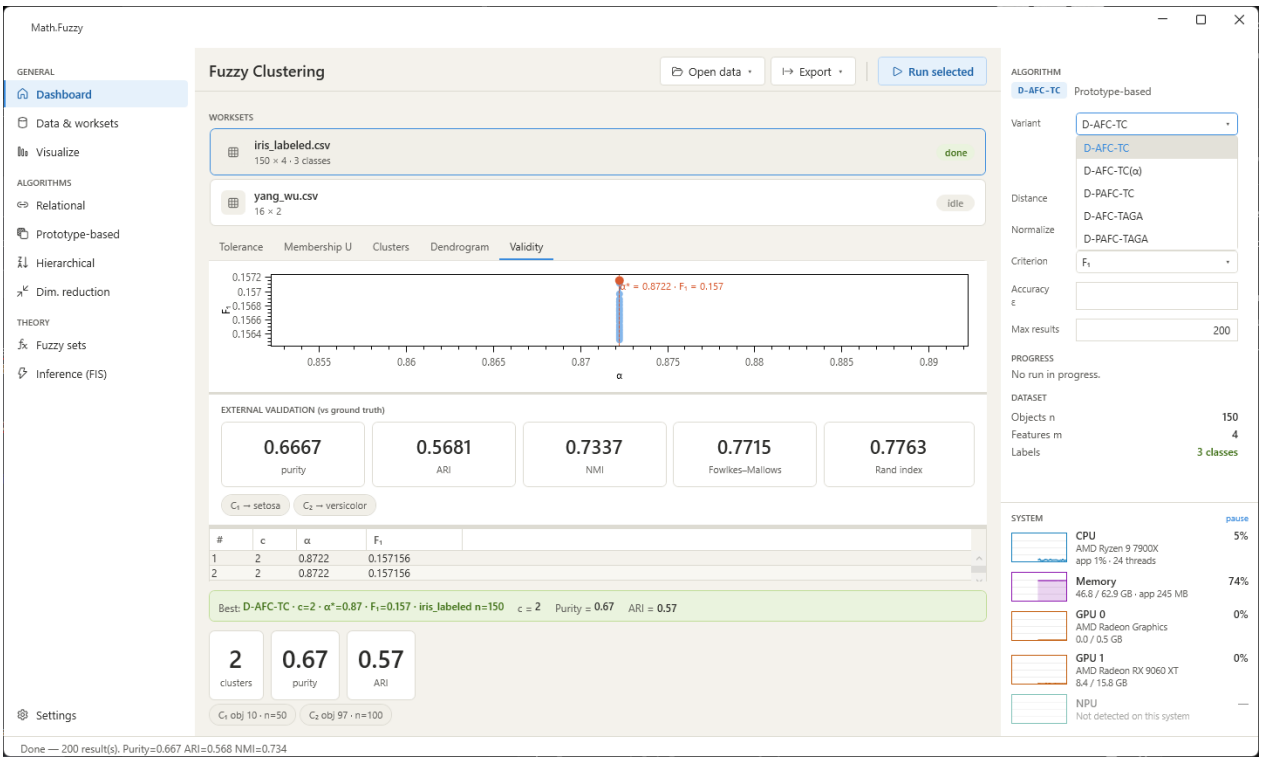


Figure 4: The prototype-based category: the *Variant* box lists its five algorithms.

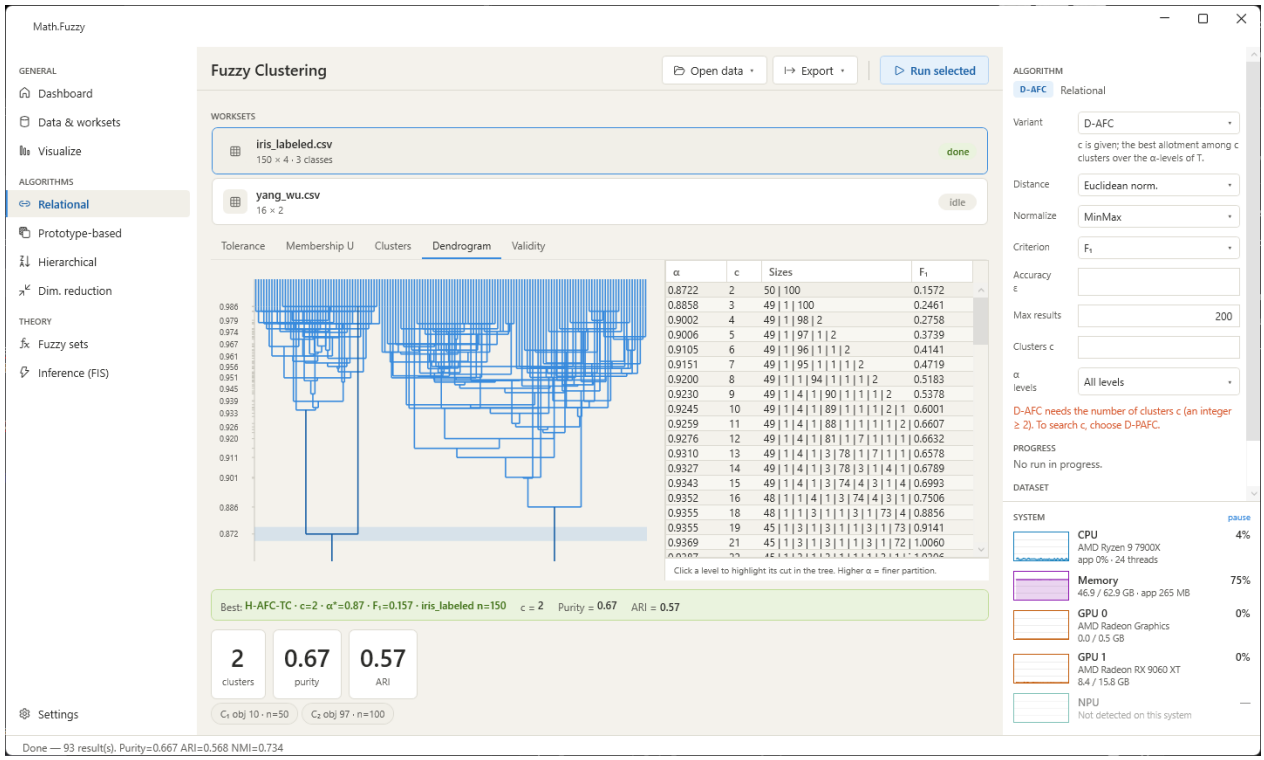


Figure 5: The relational category with D-AFC selected; the fields *Clusters* c and *levels* belong to this algorithm. Until c is entered the panel reports the missing parameter in red.

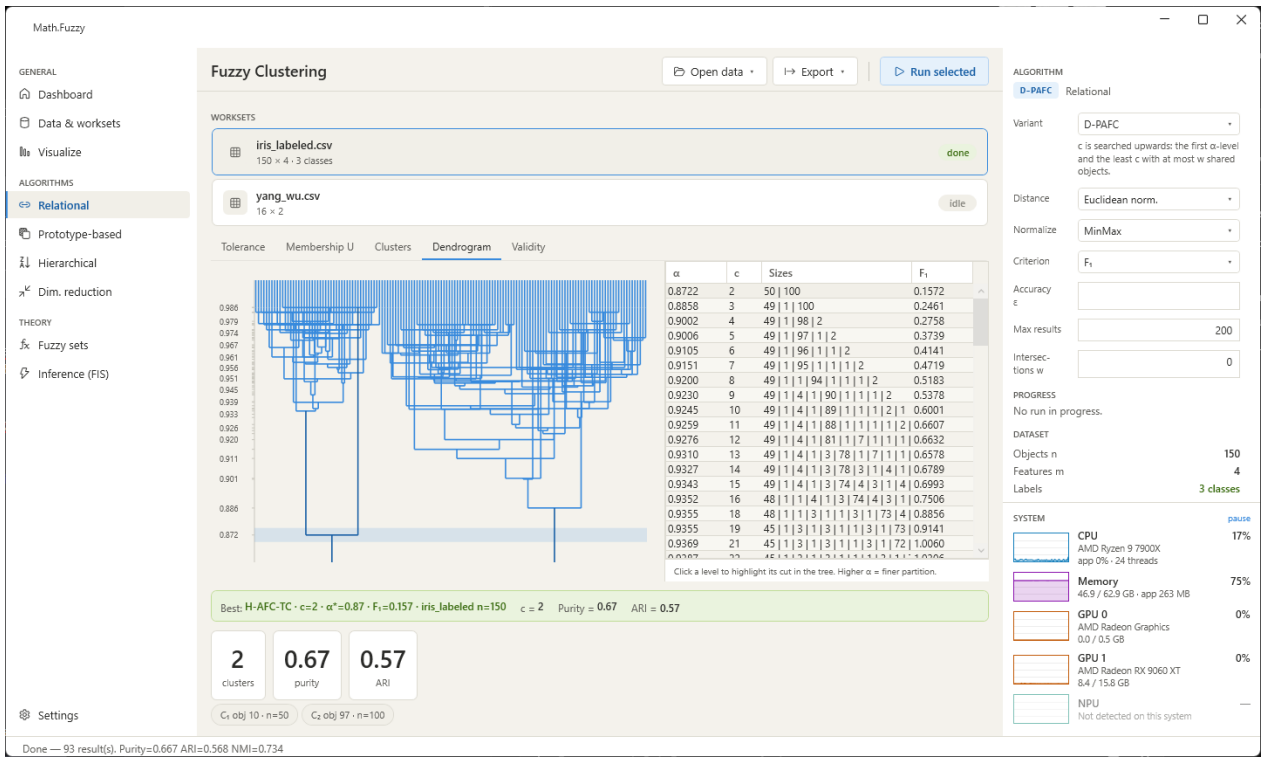


Figure 6: D-PAFC does not need c; instead the field *Intersections* w sets how many objects may belong to more than one cluster.

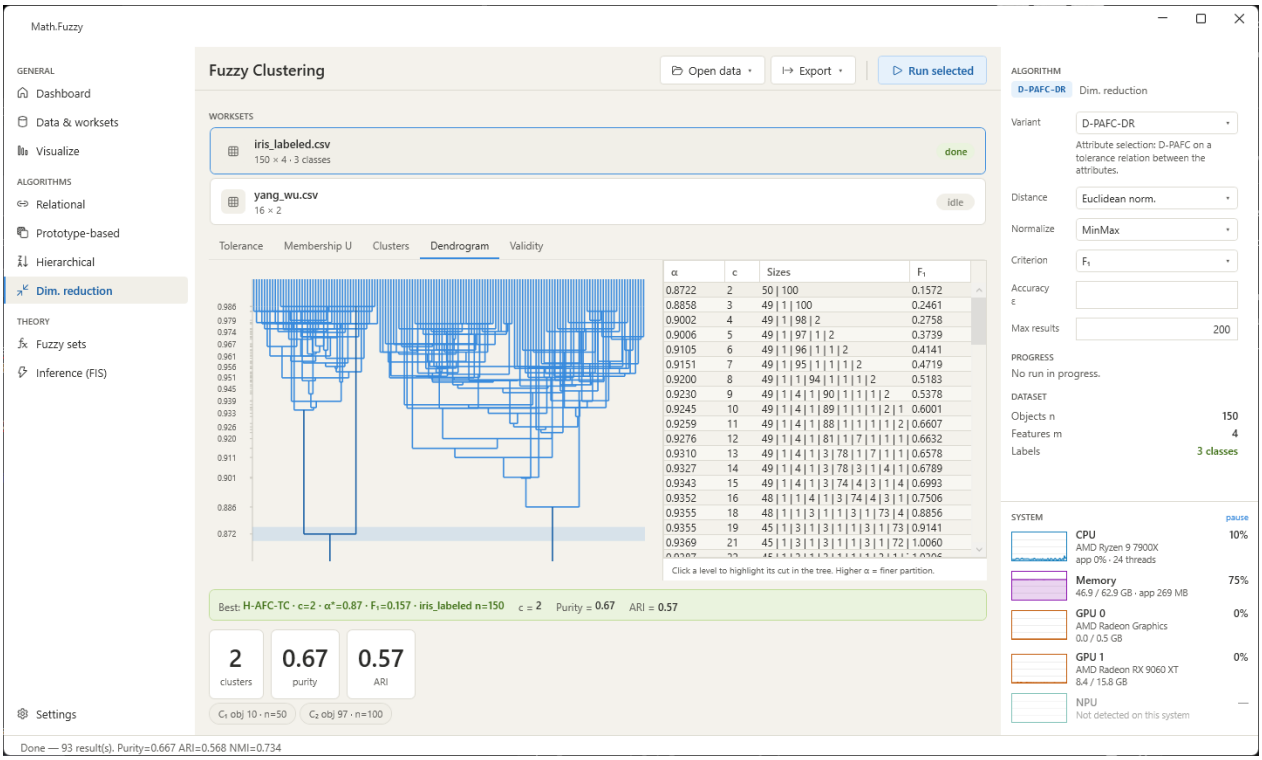


Figure 7: The dimensionality-reduction category.

5.1 Parameters

The upper part of the panel is the same for every algorithm:

- **Distance.** How the dissimilarity of two objects is measured on the normalised attributes: the normalised Hamming distance, the normalised Euclidean distance, its square, the Manhattan, Chebyshev, Minkowski (with its order p) and cosine distances. The tolerance is $T_{ij} = 1 - d(x_i, x_j)$.
- **Normalize.** *MinMax* maps every attribute onto $[0, 1]$ (unitarisation); *Max* divides every attribute by its largest value (the normalisation of the original experiments); *None* leaves the data as they are.
- **Criterion.** F_1 or F_2, the two criteria by which allotments of the same level are compared; F_1 penalises the number of clusters.
- **Accuracy** . When set, every entry of T is rounded to the nearest multiple of before the levels are built. This merges levels that differ by less than and enlarges the kernels of the clusters. Leave the field empty for full precision. Typical values are 0.00001 to 0.01.
- **Max results.** How many allotments are kept in the result list.

The lower part depends on the algorithm:

- **Clusters c** (D-AFC): the number of clusters, an integer of at least 2.
- **levels** (D-AFC): *All levels* searches every distinct level of T; *Manual* searches only the level nearest to the *** threshold** entered below.
- *** threshold** (D-AFC-TC(α)): the lowest admissible level.
- **Intersections w** (D-PAFC): the number of objects allowed to lie in more than one cluster; 0 asks for fully separated clusters.
- **TAGA operator** (TAGA variants): the aggregation operator of the transitive approximation.

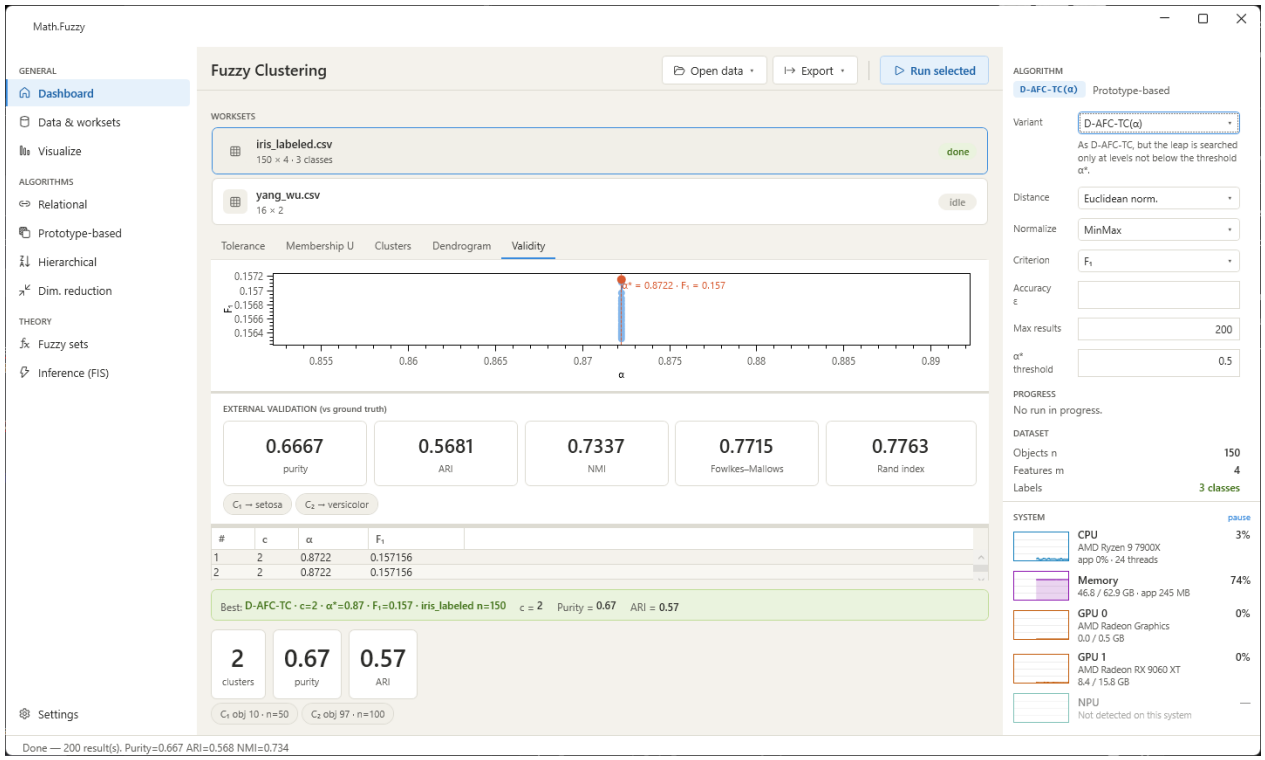


Figure 8: D-AFC-TC( ) with its field * *threshold*.

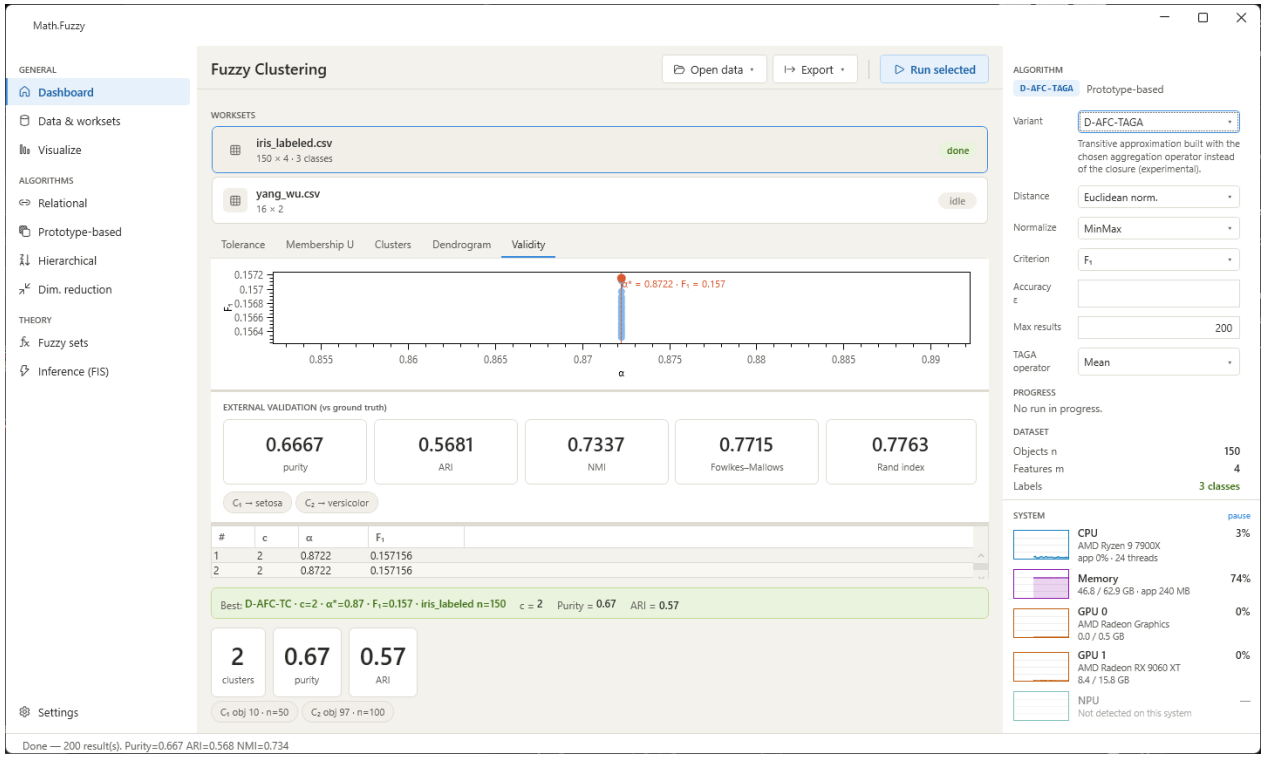


Figure 9: A TAGA variant with its field *TAGA operator*.

A parameter that is missing or out of range is reported in red under the fields, and *Run selected* does nothing until it is corrected.

Reproducing the experiments of the book

The original experiments on the Iris data use the squared Euclidean distance, normalisation *Max* and the accuracy = 0.00001. With D-AFC and c = 3 this setting gives the allotment of the monograph with six misclassified objects.

6 Running an algorithm

Select the workset, configure the panel and press *Run selected*. The workset shows *running…*, the header shows the elapsed time and *Stop*, and the progress bar reports the scan of the levels. Most runs on the sample data take well under a second; the relational algorithms on a few hundred objects can take minutes, because they enumerate combinations of centres on every level. *Stop* interrupts a run at the next level; the results of the previous run stay on the tabs.

When the run finishes, the status bar prints the number of allotments found and — for labelled data — purity, the adjusted Rand index and the normalised mutual information of the best one. If no allotment satisfies the parameters, the status bar says so; typical remedies are a different distance, *All levels* instead of a manual level, or a larger w.

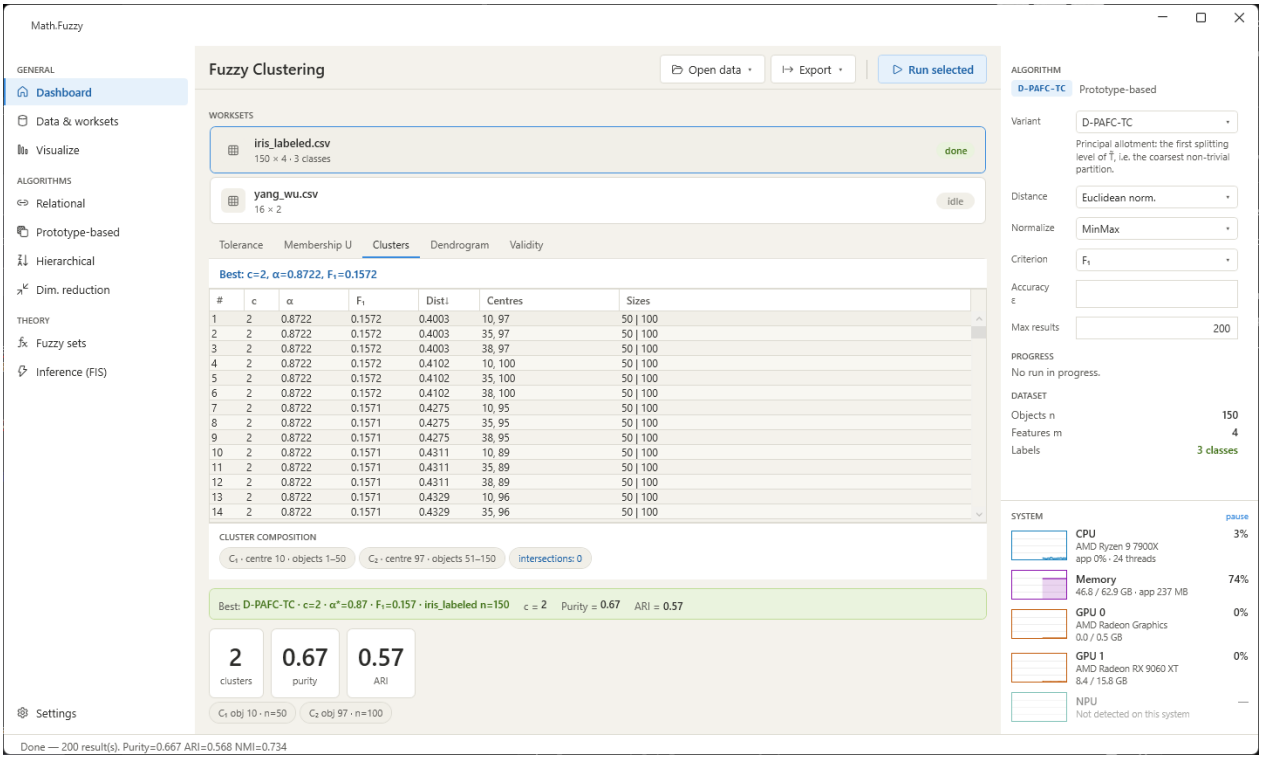


Figure 10: A D-PAFC-TC run on the Iris data has just finished.

7 Reading the results

7.1 Clusters

The table lists the allotments found, best first: the rank, the number of clusters, the level , the value of the criterion, the prototype distance used to break ties (closure-based algorithms only), the centres (typical points, 1-based object numbers) and the sizes of the clusters. Selecting a row shows its composition below the table: every cluster with its centre and the ranges of its objects, and the number of objects lying in more than one cluster.

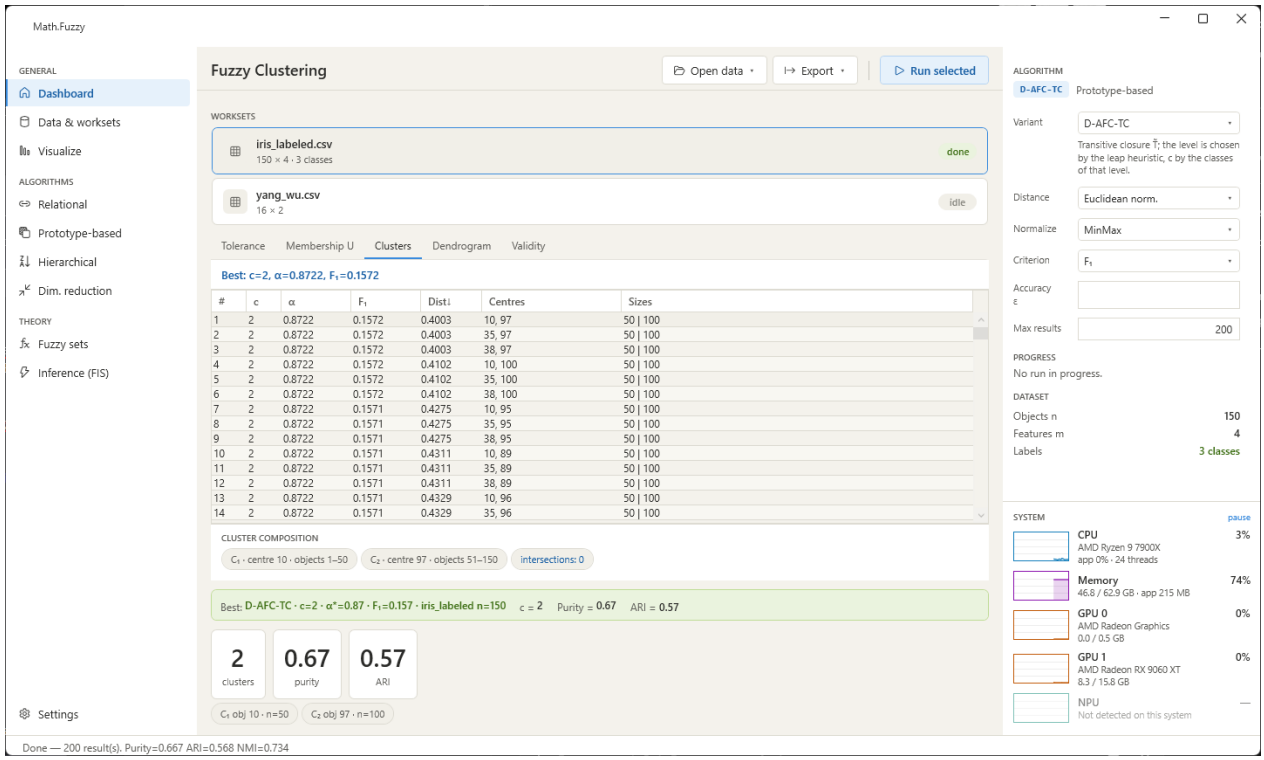


Figure 11: The list of allotments after a D-AFC-TC run.

7.2 Membership U

The membership matrix of the selected allotment as a heat map: one row per object, one column per cluster, the colour giving the membership degree μ. The *threshold* slider hides memberships below a value, *sort by cluster* orders the objects by their cluster. Clicking a row shows the memberships of that object on the right.

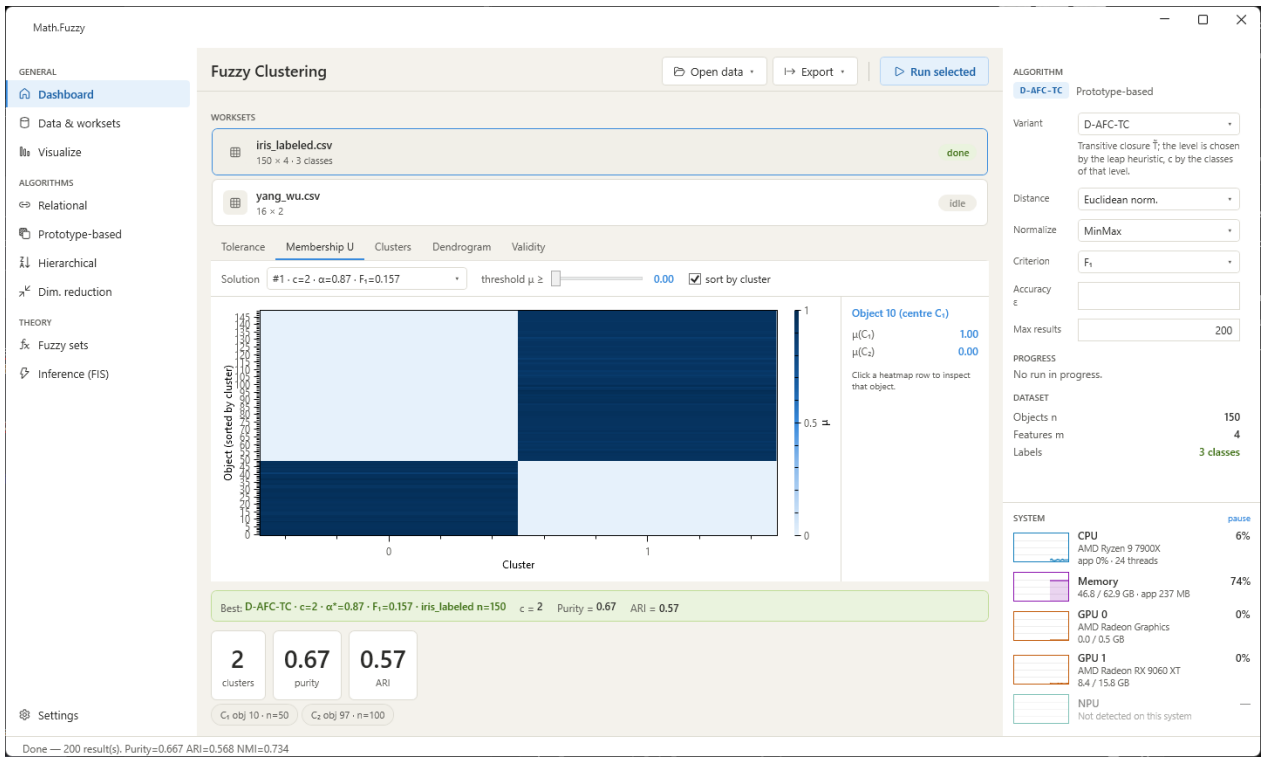


Figure 12: The membership matrix.

7.3 Tolerance

The tolerance matrix T built with the parameters of the last run, and — with the button *T (closure)* — its max-min transitive closure. The number of distinct levels is shown at the right; the *-slice* slider blanks every entry below the chosen level, which shows the crisp relation T_α at that level. Objects are in file order.

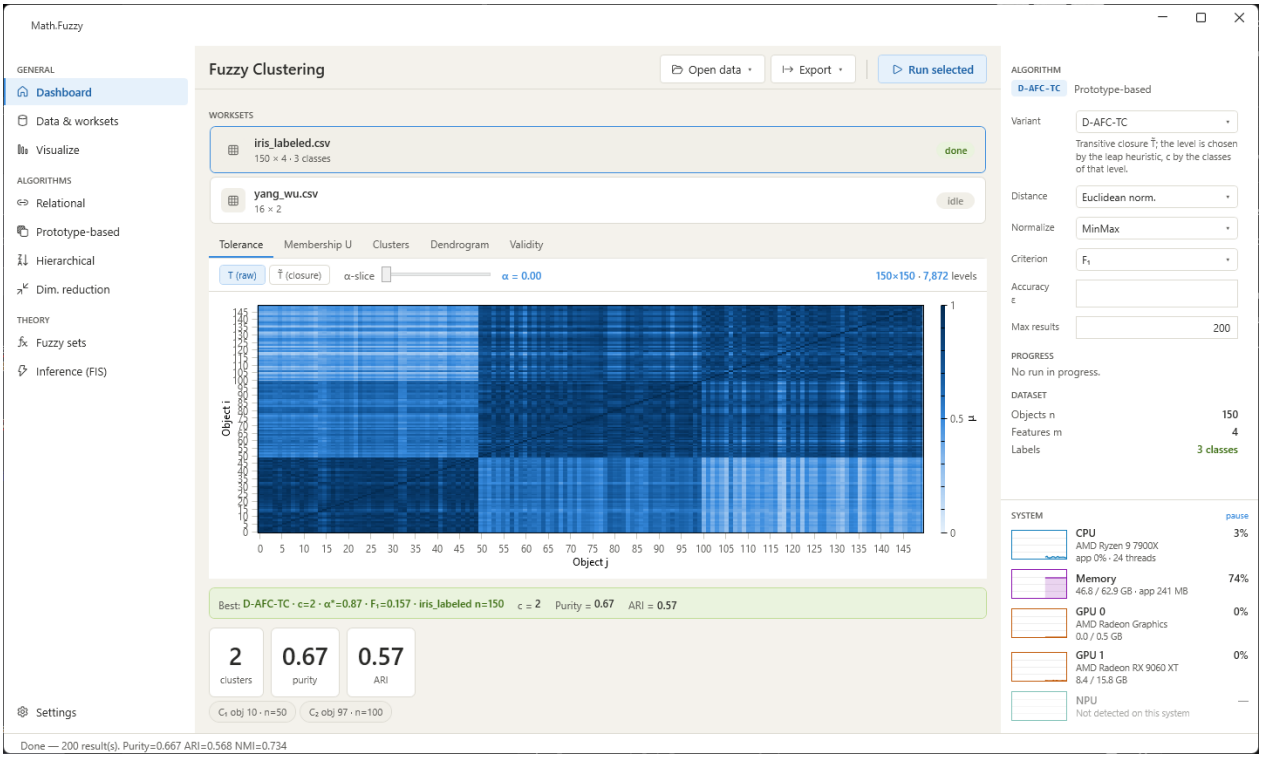


Figure 13: The tolerance matrix T of the Iris data.

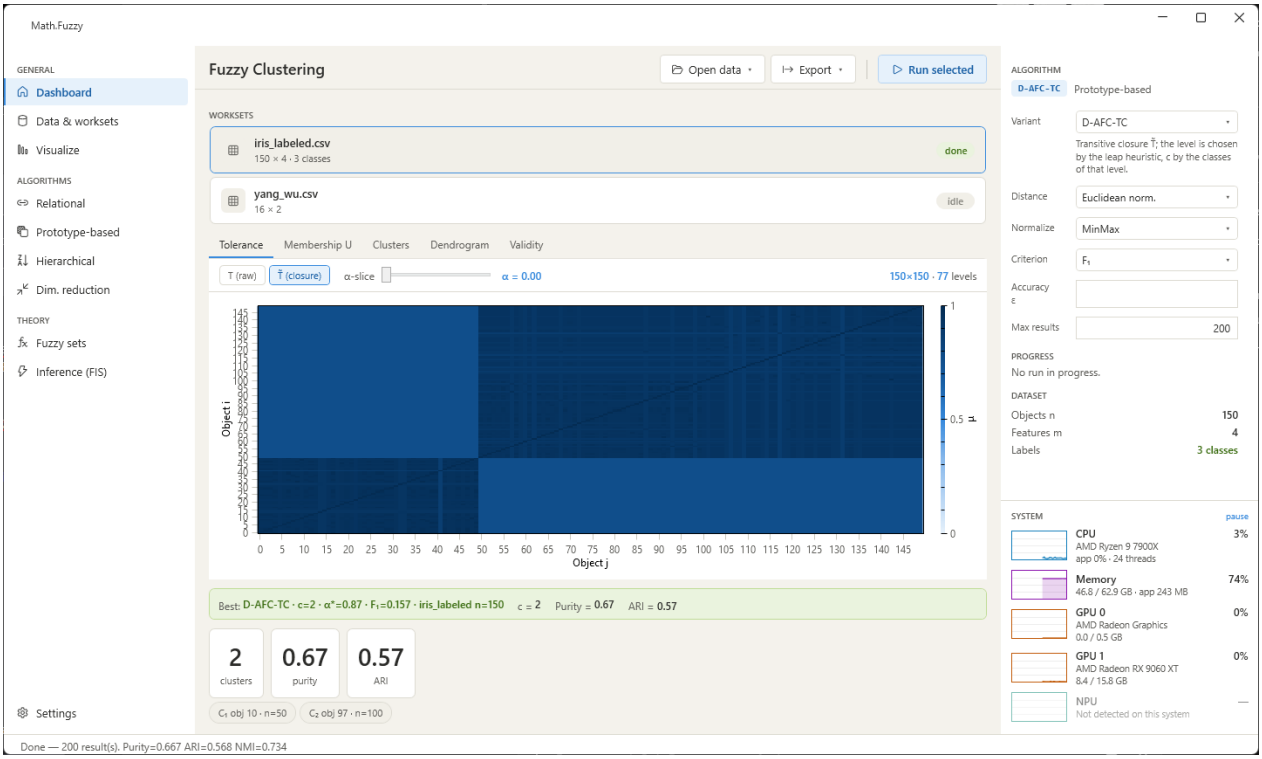


Figure 14: Its transitive closure T: the block structure of the equivalence classes.

7.4 Dendrogram

After a run of H-AFC-TC the tab shows the hierarchy of allotments: the levels of $\widetilde{T}$ on the vertical axis, the finest partition at the top, merges at decreasing going down. The table on the right lists every level with its number of clusters, the cluster sizes and the criterion; clicking a level highlights its cut in the tree. For other algorithms the tab is empty.

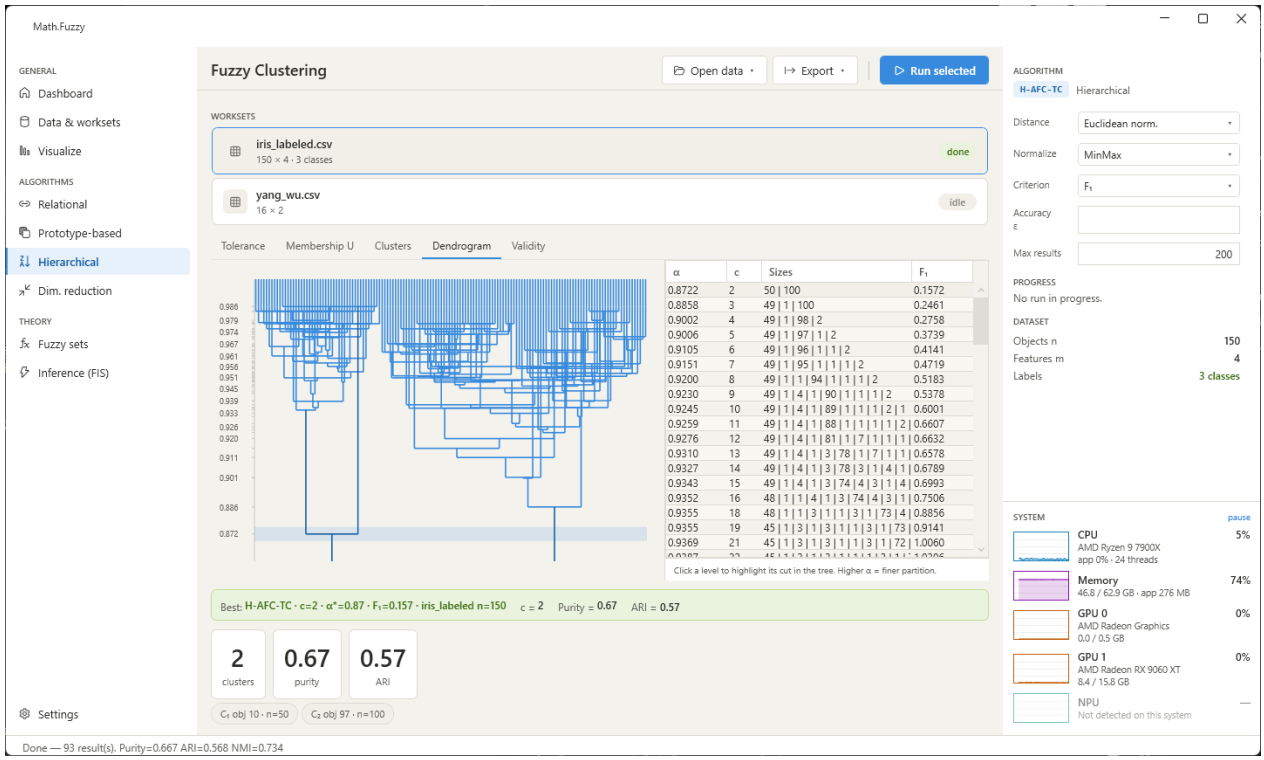


Figure 15: The dendrogram of the Iris data.

7.5 Validity

The upper plot shows the criterion of every allotment against its level, with the best one marked. For labelled data the block *External validation* gives purity, the adjusted Rand index, the normalised mutual information, the Fowlkes–Mallows index and the Rand index of the best allotment, and the majority class of every cluster.

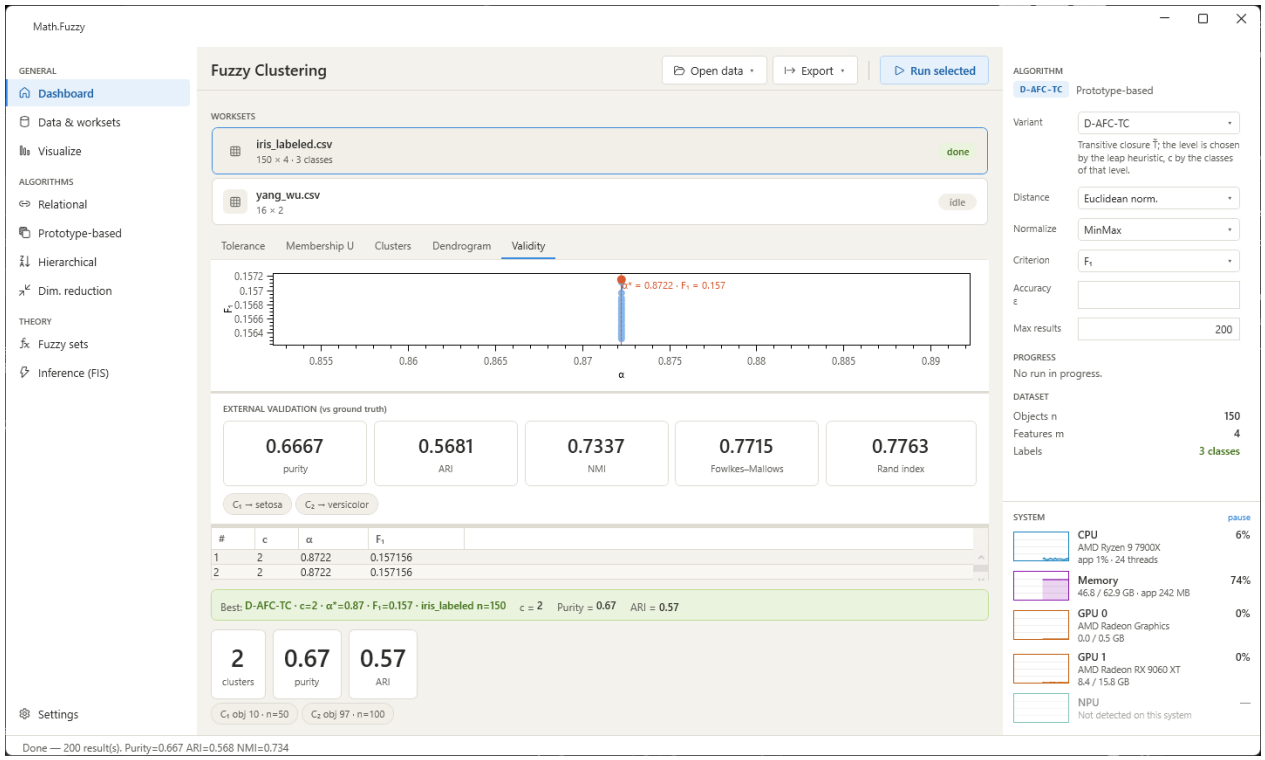


Figure 16: The validity tab after a D-AFC-TC run on the Iris data.

8 Visualize

The scatter plot shows the active workset in two chosen attributes. *Color by* switches between the clusters of the best allotment of the last run — with the typical points as black diamonds and objects of several clusters ringed — and the true classes of the file. The caption under the toolbar says which run and which data set the colours refer to.



Figure 17: The scatter plot coloured by the true classes.

9 Export

Export → Results as JSON… writes every allotment of the last run with its level, criterion, centres, clusters and the validity indices; the structure is the same as the JSON output of the command-line tool. *Export → Results as CSV…* writes one row per allotment for spreadsheets.

10 Theory pages

10.1 Fuzzy sets

A small laboratory for membership functions. *Shapes* draws one fuzzy set of a chosen shape (triangular, trapezoidal, Gaussian and others); the parameters can be typed or the square handles dragged on the plot, and the panel reports the height, support, core, -cut, crossover points, cardinality and centroid. *Operations* combines two sets with a t-norm, an s-norm or a complement, with every operator of the chosen family, and defuzzifies the result by five methods. *From cluster* shows how a fuzzy cluster of the last run becomes a term of the antecedent of a rule.

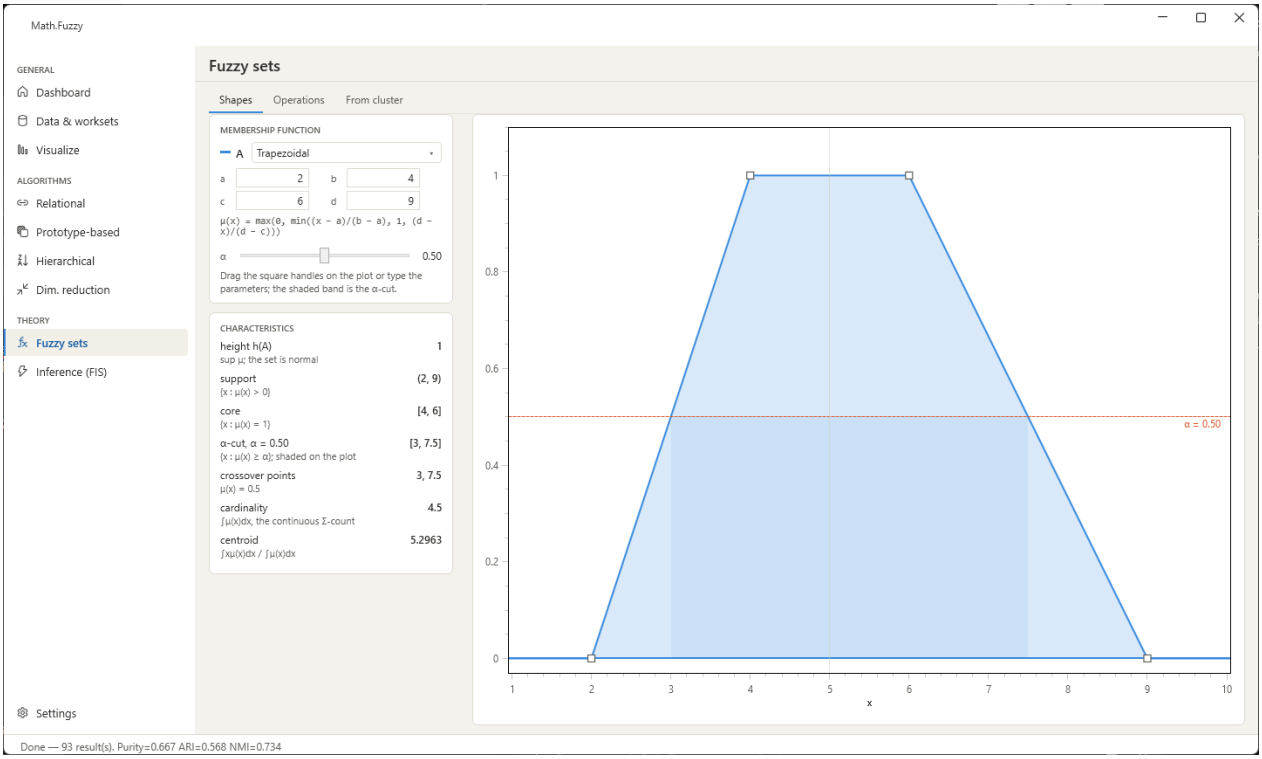


Figure 18: The Shapes tab.

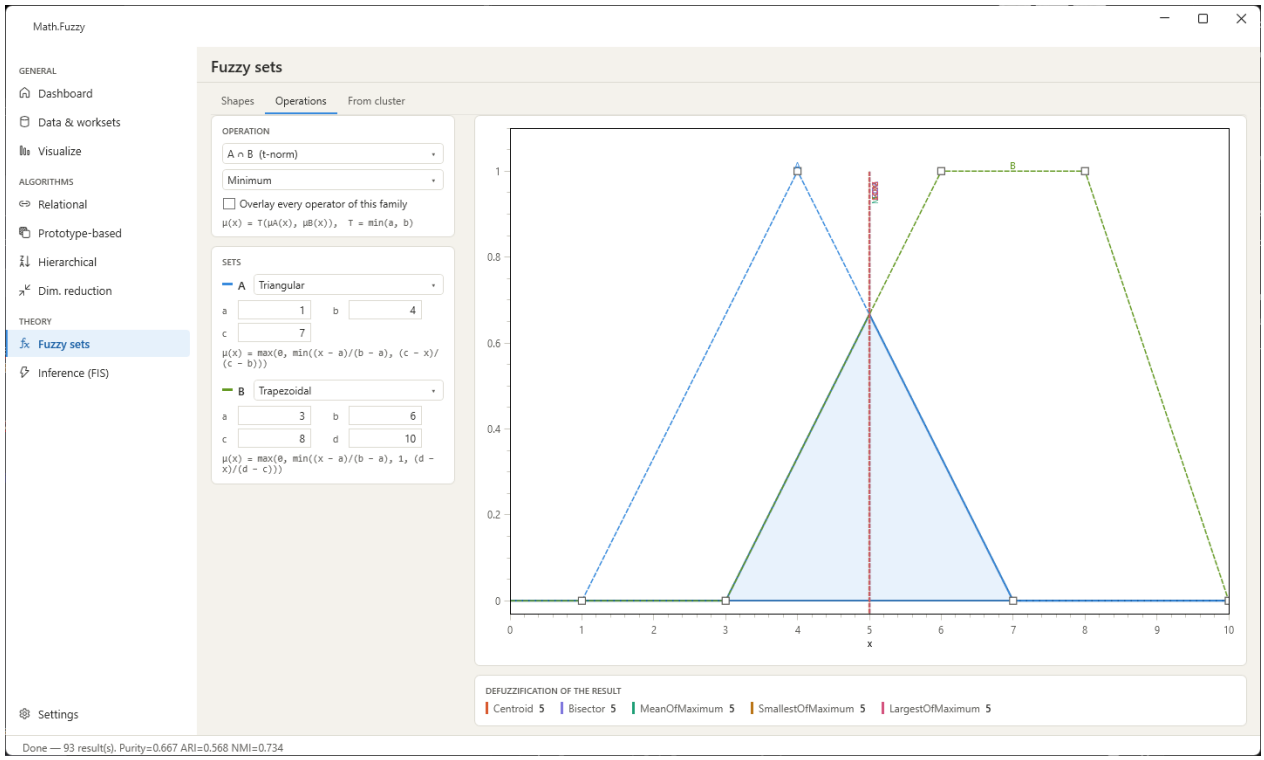


Figure 19: The Operations tab: the intersection of two sets.

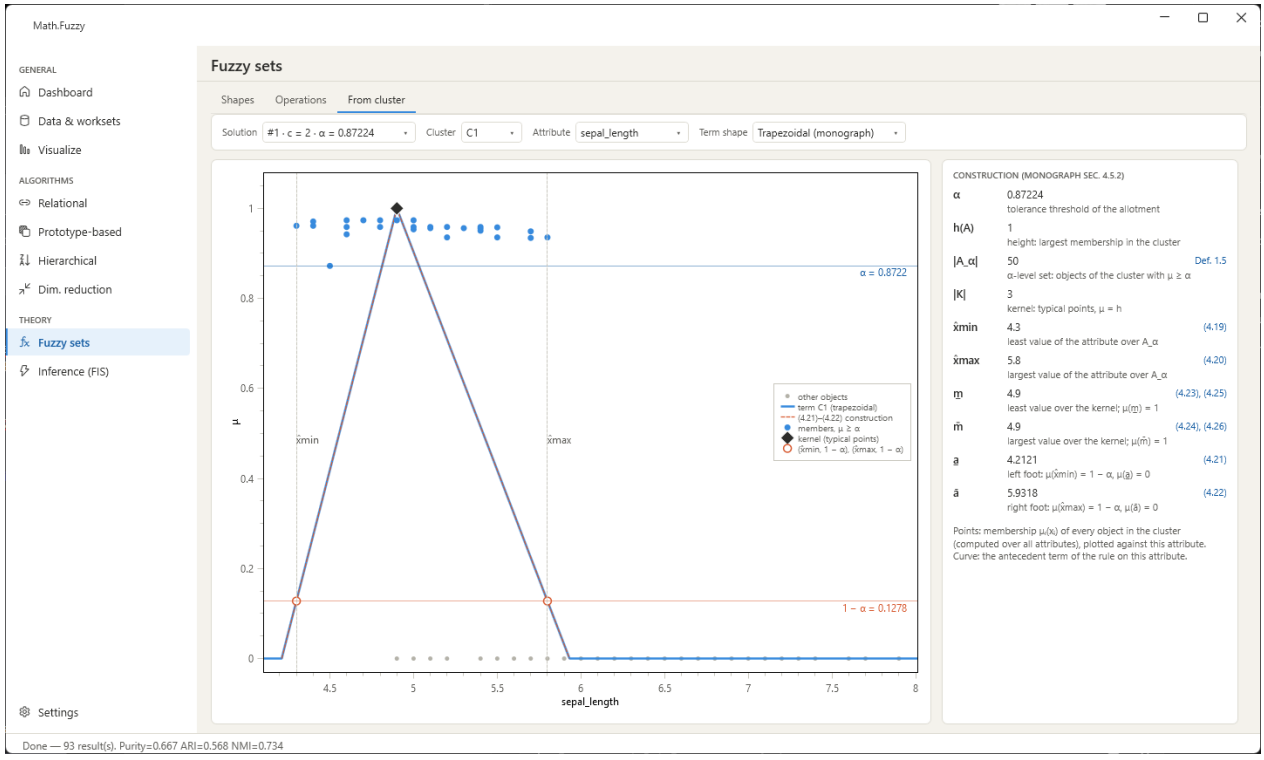


Figure 20: The From cluster tab: the antecedent term built from cluster C1 of a D-AFC-TC solution, with the construction values on the right.

10.2 Inference (FIS)

The inference page generates a fuzzy rule base from a clustering solution and lets you experiment with it.

- **Solution** picks an allotment of the last run; **Rule base** the kind of system (a Mamdani system with

one output per cluster, a Takagi–Sugeno system or a classifier whose consequents are the class labels); **MF shape** the shape of the terms. *Regenerate* builds the rules.

- **Rules.** One rule per cluster. Each rule can be switched off, its terms and weight changed; the bar at the right shows its firing strength for the current input.
- **Membership functions.** The terms of one variable; parameters can be typed or dragged.
- **Input vector.** Sliders for every attribute; *Load* copies an object of the data set into them.
- **Output.** The decision, the output values and the aggregated output set with its defuzzified value.
- **Evaluation on the workset.** The accuracy of the rule base on all objects and its confusion matrix.

The rule base can be exported as JSON or as a MATLAB Fuzzy Logic Toolbox file (`.fis`) and imported back.

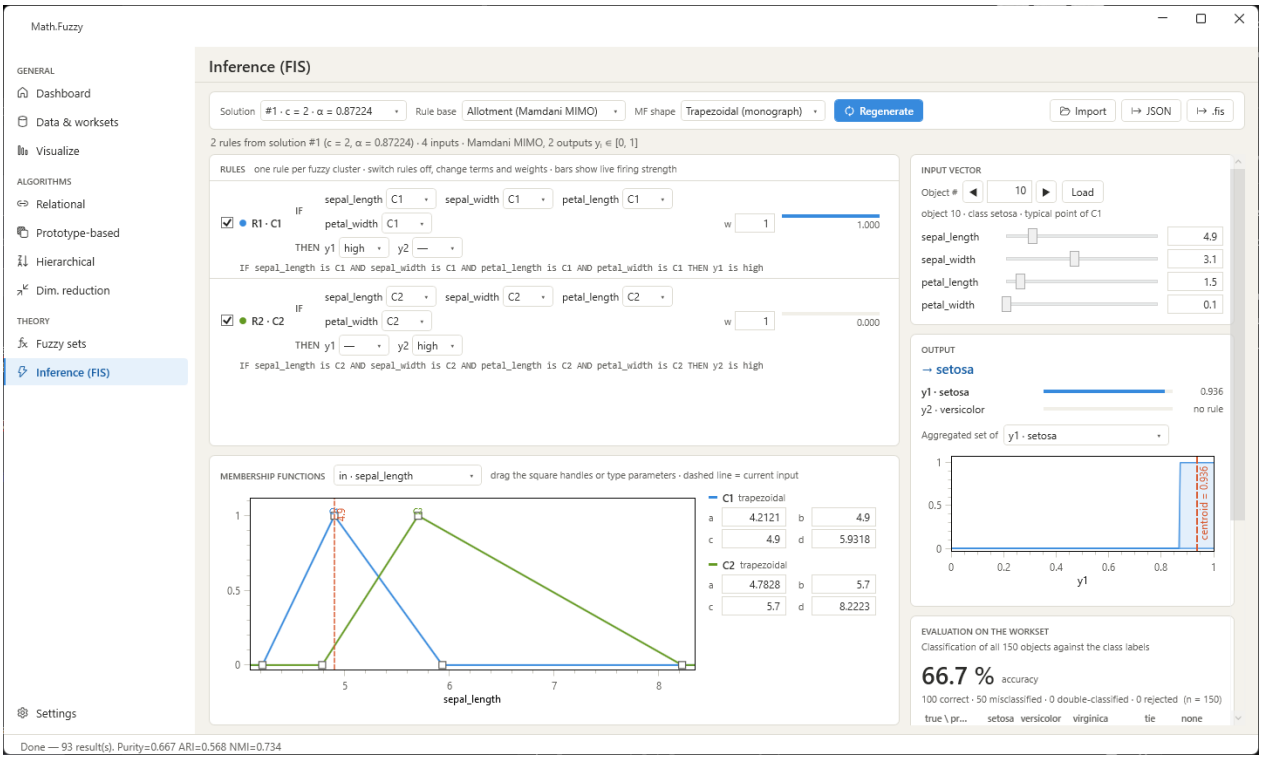


Figure 21: The inference page with a rule base generated from a D-AFC-TC solution.

11 Settings

The page sets the defaults for the number of results kept and the minimal size of a cluster, and shows the versions of the components.

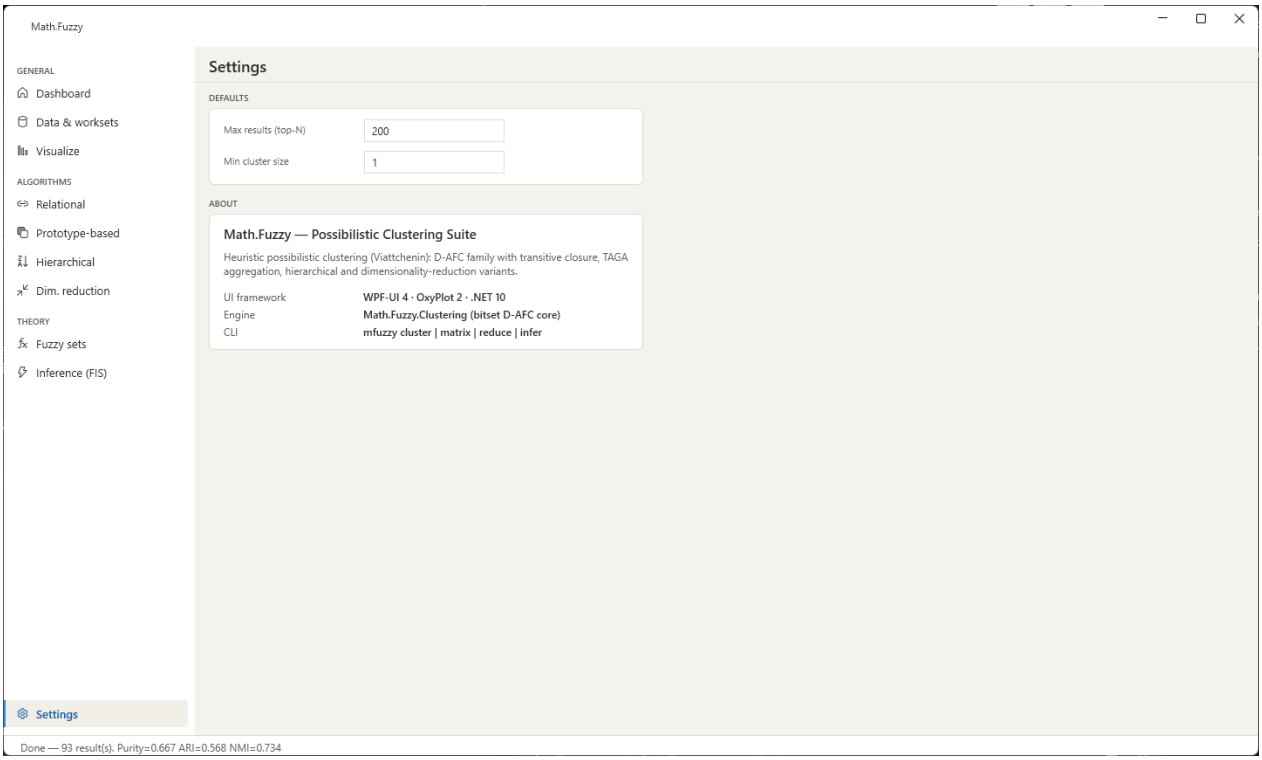


Figure 22: The Settings page.

12 Troubleshooting

- *Windows refuses to start the application.* SmartScreen: choose *More info* → *Run anyway*. The checksum of the archive can be verified against the `.sha256` file published with the release.
- *"D-AFC needs the number of clusters c."* Enter an integer of at least 2 in *Clusters c*, or choose D-PAFC to have the number searched.
- *"No valid clustering found."* No allotment satisfies the parameters. With a manual level try *All levels*; with D-PAFC increase w; try another distance.
- *A run takes minutes.* The relational algorithms enumerate centre combinations on every level; a few hundred objects with $c = 3$ can take several minutes, and partially separated clusters ($w > 0$) rescan all levels. The closure-based algorithms are fast. *Stop* interrupts a run.
- *A workset loaded as a matrix cannot be run.* The algorithms of this version take attribute data; a matrix workset can be inspected on the *Tolerance* tab.
- *The TAGA variants give unexpected results.* The transitive approximation is experimental in this version and its results are not reported in the book.
- *The results of a previous run are still shown after an error.* They belong to the previous run; the status bar reports the error of the new one.

13 The command-line tool

Everything the workbench computes can be reproduced from a terminal with the tool `mfuzzy` of the same repository, for example

```
mfuzzy cluster --algorithm dafc --input datasets/iris_labeled.csv --clusters 3 ^
    --distance squared-euclidean --normalize max --accuracy 0.00001
```

The options carry the names of the panel fields (`--distance`, `--normalize`, `--criterion`, `--accuracy`, `--clusters`, `--alpha`, `--intersections`, `--taga`), and `--format json` writes the same JSON as the *Export* command. The README of the repository documents the tool, the .NET library and the MATLAB toolbox.